

COMP1002

DATA STRUCTURES AND

ALGORITHMS

LECTURE 4: LINKED LISTS



Curtin University

Discipline of Computing

Last updated: [March 16, 2020](#)

Copyright Warning

COMMONWEALTH OF AUSTRALIA

Copyright Regulation 1969

WARNING

This material has been copied and communicated to you by or on behalf of
Curtin University of Technology pursuant to Part VB of the Copyright
Act 1968 (the Act)

The material in this communication may be subject to copyright under the
Act. Any further copying or communication of this material by you may be
the subject of copyright protection under the Act.

Do not remove this notice

This Week

- Linked Lists
 - Simple linked lists
- Variants
 - Double-ended
 - Doubly-linked
 - Sorted linked lists
- Time complexity analysis
 - Big-O notation
- Serialization
- Iterators

Arrays Aren't Everything

- As a data structure, arrays have their disadvantages
 - **Fixed size** – cannot grow/shrink to fit num elements
 - If we don't know how many elements in advance, we'll need to overestimate the array's length, and thus waste space
 - **Insertion** – inserting an element between two existing elements will require shuffling later elements up by one
 - Similar issues plague **removal** from the front/middle (e.g., queues)
 - **Contiguous** – all elements must fit in a contiguous block
 - This is sometimes a disadvantage – it's a bit inflexible
- We'd like a data structure that doesn't suffer from these
 - ... but still does a similar job to arrays (storing list of data)
 - Fortunately, it's already been invented: the linked list


Linked Lists

- The idea behind a linked list is that each element stores:

a **value** and a **reference** to the *next* element

- The list is essentially a chain of links
- The list itself only stores a reference to the *first* element
 - This is often called the 'head'
- Since each element only points to the next, it is *singly-linked*
 - Since the list only points at the head, it is *single-ended*
- More complicated variants on the basic single-ended singly-linked list
 - Double-ended, doubly-linked, etc
 - We'll examine some of these variants later

Simple Linked List



Anatomy of a Linked List

- Each element is not just data but also a reference to the next element
 - Thus we need to create a composite of `value` + `next` ref
 - Classes are ideally suited to this – make `value`, `next` class fields
 - Typically called `ListNode`, or less commonly `ListLink`
 - In pre-OO languages, use structs (C) or records (Pascal)
- Head is the only member field of `LinkedList` class
 - Points at the first `ListNode` in the list
- Last node's next points at null
 - Indicates the end of the linked list

ListNode Members

- ListNode only exists as a container for the **data** and the **next** reference
- Thus it is really only a couple of member fields with associated getters/setters (Java):
 - One member field (instance variable) for the **data** value, usually an Object so that the list can be general-purpose
 - A member **next** that is a reference to the next node in the list's chain
 - This points at null/None if the ListNode is the last node in the chain

Traversing Linked Lists

- A linked list is just a chain of connected-but-independent nodes
 - Each node could be anywhere in memory
 - Only a node's predecessor knows where to find a given node (in memory)
 - So the only way to get to the fourth node is to first get to the third node
 - which in turns requires us to have made it to the second node
 - ... and we can only get to that via the first node

Traversing Linked Lists

- Only the first node is available directly (via *head*)
 - Thus getting to a particular node requires that we traverse from the *head* (first) node, through all the *next* pointers until we reach the desired node
 - This obviously can be pretty slow
 - Compare this to an array, which can get to *any* element in *one step*
 - via a little bit of arithmetic – it works because the array elements are all stored in a contiguous block
 - But linked lists aren't contiguous, so we have to do traversing
- This is the main disadvantage of linked lists: *access time*

Linked List Methods

- isEmpty()
- insertFirst(), insertLast(), insertBefore(valueToFind)
 - Insert a new item into the list
 - Should require a data item as import, NOT a ListNode
 - ListNode is an **internal** detail of LinkedList's implementation
- removeFirst(), removeLast(),
 - Delete an item from the list
- peekFirst(), peekLast(),
 - Return the **data value** of an item in the list (not ListNode)
- find(<dataType> valueToFind)
 - Search whether a given data value exists in the list


remove(valueToFind)
(less common)
peek(valueToFind)

Inserting a ListNode

- There are four possible scenarios for node insertion:
 - The list is empty and we are inserting the first node
 - We are inserting before the head node
 - We are inserting after the tail of the list
 - We are inserting somewhere in the middle of the list
- In code, some of these cases turn out to be the same thing, but when designing a solution you should *always* be thinking through all the possibilities


Inserting First ListNode

- Initial: empty list
(*i.e.*, head points at null/None)




- Step 1: Create new ListNode with (say) value of 3.8

<pre>ListNode newNd; newNd = new ListNode(3.8);</pre>	<pre># Python newNd = ListNode(3.8)</pre>
---	---




- Step 2: Point head at the newNd to make it the first node




Inserting ListNode Before Head

- Initial: a few nodes in the list




- Step 1: Create new ListNode with (say) value of 15.2

```
ListNode newNd;  
newNd = new ListNode(15.2);  
# Python  
newNd = new ListNode(15.2)
```




- Step 2&3: Fix links to make newNd first

```
newNd.next = head;  
head = newNd;
```




Inserting ListNode After Tail


- Initial: a few nodes in the list



- Step 1: Create new ListNode with (say) value of 100




- Step 2: Make last node in the chain point at the newNd




Inserting ListNode In The Middle


- Initial: a few nodes in the list



- Step 1: Create new ListNode with (say) value of 10.3



- Step 2: 'Fix up' links to make it fit correctly in the list




Removing a ListNode

- As with insertion, there are four cases:
 - Removing the head ListNode
 - Removing the tail ListNode
 - Removing a ListNode in the middle of the list
 - Removing the sole remaining ListNode in the list
- And as with insertion, some of these cases end up being the same in the code


Removing Head ListNode

- Initial: a few nodes in the list



```
head = head.next;
```


- Step 1: Point head at second node




- First node is now no longer part of the list and its memory usage will be garbage-collected by Java (in C/C++, you would have to explicitly free it)
 - The fact that Node 15.2 still points at the second node is not relevant: nothing points *at* Node 15.2, therefore it will be garbage-collected

Removing Tail ListNode

- Initial: a few nodes in the list



- Step 1: Point second-last node at null.
Last node is now no longer part of the list




```
ListNode currNd, prevNd;
currNd = head;
prevNd = null;

// Traverse to last node
while (currNd.next != null) {
    prevNd = currNd;
    currNd = currNd.next;
}
prevNd.next = null;
```

```
# Python
currNd = head
prevNd = null
// Traverse to last node
while (currNd.next)
    prevNd = currNd
    currNd = currNd.next
prevNd.next = None
```


Removing ListNode In The Middle

- Initial: a few nodes in the list;




- Node 10.3 is to be removed

-
- Step 1: Have predecessor node point at Node 10.3's next node.
Node 10.3 is now no longer part of the list




Removing Last Remaining ListNode

- Initial: one node remains in the list



-
- Step 1: Point head at null



- Node is now no longer part of list

```
// Same as removing first node  
head = head.next;
```

ListNode - Pseudo-code

```
Class DSAListNode
Class fields : value (Object), next (DSAListNode) // Could make these public members
                                                    // and avoid the need for getters/setters since
                                                    // ListNode is only used INSIDE LinkedList

Alternate constructor IMPORT inValue (Object)      // Only need one constructor
    value ← inValue                                // No need for validation
    next ← NULL

ACCESSOR getValue IMPORT none EXPORT value

MUTATOR setValue IMPORT inValue (Object) EXPORT none
    value ← inValue

ACCESSOR getNext IMPORT none EXPORT next

MUTATOR setNext IMPORT newNext (DSAListNode) EXPORT none
    next ← newNext
```

Simple LinkedList - Pseudo-code

```
Class DSALinkedList
Class fields : head (DSAListNode)
```

```
Default constructor
    head ← NULL
```

```
MUTATOR insertFirst IMPORT newValue (Object) EXPORT none
    newNd ← allocate DSAListNode(newValue)
    IF isEmpty() THEN
        head ← newNd
    ELSE
        newNd.setNext(head)
        head ← newNd
    ENDIF
```

```
MUTATOR insertLast IMPORT newValue (Object) EXPORT none
    newNd ← allocate DSAListNode(newValue)
    IF isEmpty() THEN
        head ← newNd
    ELSE
        currNd ← head
        WHILE currNd.getNext() <> NULL DO           // Traverse to last node <> is "not equal"
            currNd ← currNd.getNext()
        ENDWHILE
        currNd.setNext(newNd)
    ENDIF
```

<continued next slide>

Simple LinkedList - Pseudo-code

```
ACCESSOR isEmpty IMPORT none EXPORT empty (boolean)
empty ← (head = NULL)
```

```
ACCESSOR peekFirst IMPORT none EXPORT nodeValue (Object)
  IF isEmpty() THEN
    abort
  ELSE
    nodeValue ← head.getValue()
  ENDIF
```

```
ACCESSOR peekLast IMPORT none EXPORT nodeValue (Object)
  IF isEmpty() THEN
    abort
  ELSE
    currNd ← head
    WHILE currNd.getNext() <> NULL DO // Traverse to last node
      currNd ← currNd.getNext()
    ENDWHILE
    nodeValue ← currNd.getValue()
  ENDIF
```

<continued next slide>

Simple LinkedList - Pseudo-code

```
MUTATOR removeFirst IMPORT none EXPORT nodeValue (Object)
```

```
  IF isEmpty() THEN
```

```
    abort
```

```
  ELSE
```

```
    nodeValue ← head.getValue()
```

```
    head ← head.getNext()
```

```
  ENDIF
```

```
MUTATOR removeLast IMPORT none EXPORT nodeValue (Object)
```

```
  IF isEmpty() THEN
```

```
    abort
```

```
  ELSEIF head.getNext() = NULL THEN
```

```
    nodeValue ← head.getValue()
```

```
    head ← NULL
```

```
  ELSE
```

```
    prevNd ← NULL
```

```
    currNd ← head
```

```
    WHILE currNd.getNext() <> NULL DO
```

```
      prevNd ← currNd
```

```
      currNd ← currNd.getNext()
```

```
    ENDWHILE
```

```
    prevNd.setNext(NULL)
```

```
    nodeValue ← currNd.getValue()
```

```
  ENDIF
```

```
// Traverse to last node
```

```
// We need to get the second-last node
```

```
// in order to 'drop' the last node from the list
```

```
// Remove currNd from list
```

```
// Return value of node that is being removed
```

Private Inner classes (Java)

```
public class DSALinkedList
{
    private class DSAListNode
    {    // Private inner class
        private Object m_value;
        private DSAListNode m_next;
        //could make the classfields public as they can
        //only be seen inside DSALinkedList
        public DSAListNode(Object inValue)
        {
            m_value = inValue;
            m_next = null;
        }
        // normal accessors and mutators if required
    }    // end private inner class DSAListNode
    // class DSALinkedList continues
```

Simple LinkedList / ListNode

- NOTE WELL:
 - In the practicals, you are supposed to implement a *double-ended* linked list by adjusting this pseudocode
 - (this pseudocode is just a single-ended linked list)
 - We will be discussing double-ended lists in a moment

Linked Lists vs Arrays

- Advantages of linked lists over arrays
 - ☑ Can grow and shrink with data - no space wasted
 - ☑ No limits on size (besides memory available)
 - ☑ Easy to insert items at any part of the list
 - Arrays are only easy to add items to the end
- Disadvantages
 - ☒ Time to access any given element: $N/2$ steps (on average)
 - Where N is the number of nodes in the list
 - Array is direct access: only takes 1 step
 - ☒ `next` pointer is an overhead in terms of memory usage
 - Adds 1 word (memory address) extra per element in addition to actual data

Order of Complexity

- let's analyse LinkedList find() in Big-O notation
 - Best case: the head is the node to find
 - This is one step, so $O(1)$
 - Worst case: last node is the match, which is N node hops
 - $O(N)$ in Big-O notation
 - Each hop involves multiple CPU instructions, but we aren't concerned with these details, so we don't talk about $O(5N)$
 - Average case: On average, we must go halfway: $N/2$ hops
 - $O(N)$ – again, the constant multiplying factor of $\frac{1}{2}$ is irrelevant
 - We are mostly interested in the average and worst cases


Double-Ended Linked Lists

- We often want to work at the start *and* end of a list
 - e.g., queue adds to end, takes from front
- Simple list is slowest when doing anything at the end
 - Takes $O(N)$ to reach the end – we'd like to do better!
- **Solution:** add another pointer to point at the last node
 - Call it tail
 - Put it with head as a member field of the LinkedList
 - That way we can access both head and tail in $O(1)$
- No disadvantage to this
 - Only adding one pointer to the whole linked list
 - Some code gets more complex, but other code gets easier

Double-Ended vs Single-Ended Lists

- Adding a tail pointer changes how many of the LinkedList methods work
 - insertLast() obviously becomes as easy/fast as insertFirst()
 - Same with peekLast()
 - insertFirst() and removeFirst() now have to potentially set tail as well in the case of a one-node list
 - *i.e.*, must maintain tail as well as head everywhere in the code
 - Counter-intuitively, removeLast() isn't helped much
 - We need to get at the *second*-last node to do removeLast()!


Double-Ended Linked List




Doubly-Linked Lists

- Sometimes it's useful to be able to traverse both forwards *and backwards* through a list
 - e.g., undo method in an application: in order to allow redo, must be able to go back and forth one node at a time
- Doubly-linked lists add another pointer to each ListNode: *prev*
 - **Issue:** Now overheads are *two* references per ListNode
 - That's 2 words extra *per node* (8 bytes in a 32-bit machine; 16 bytes for 64-bit)
 - Given that *ints* are only 4 bytes, it can easily *treble* space usage!
 - **Benefit:** removeLast() is easy if double-ended as well

Doubly-Linked List



Double-Ended, Doubly-Linked List



Sorted Linked Lists

- Since it's easy to insert new elements at any position in a linked list, it can also be convenient to insert them in **sorted order**
 - When adding a new element, find where it would fit (in sorted order) in the list and insert it there
- This is a variant on the Insertion Sort algorithm
 - A Doubly-Linked List makes this simple
 - no moving elements – just search backwards then update links!

Linked List Methods – for sorted list



- isEmpty()
- insertFirst(), insertLast(), **insertBefore(valueToFind)**
 - Insert a new item into the list
 - Should require a data item as import, NOT a ListNode
 - ListNode is an **internal** detail of LinkedList's implementation
- removeFirst(), removeLast(),
 - Delete an item from the list
- peekFirst(), peekLast(),
 - Return the **data value** of an item in the list (not ListNode)
- find(<dataType> valueToFind)
 - Search whether a given data value exists in the list

remove(valueToFind)
(less common)
peek(valueToFind)

Time Complexity Comparison

- Comparison of data structure speeds has many facets
 - Insertion, deletion, accessing and finding
 - Access = getting the value at a known index
 - Finding = searching through the data set to find the value
 - Best case, worst case, average case
 - Usually corresponding to front, rear and middle
 - ... but not necessarily in that order!

Complexity: Array vs Simple LinkedList

Array vs List		Insert	Delete	Access	Find
Array	At front	$O(N)$	$O(N)$	$O(1)$	$O(1)$
	At end	$O(1)$	$O(1)$	$O(1)$	$O(N)$
	In middle	$O(N)$	$O(N)$	$O(1)$	$O(N)$
Linked List (single-ended)	At front	$O(1)$	$O(1)$	$O(1)$	$O(1)$
	At end	$O(N)$	$O(N)$	$O(N)$	$O(N)$
	In middle	$O(N)$	$O(N)$	$O(N)$	$O(N)$

NOTE: Although linked lists are fast to insert/delete, it still takes $O(N)$ to traverse the list and make it to the place to insert/delete

Complexity: List Variants vs Simple List

List variants vs single-ended		Insert	Delete	Access	Find
Linked List (double-ended)	At front	$O(1)$	$O(1)$	$O(1)$	$O(1)$
	At end	$O(1)$	$O(N)$	$O(1)$	$O(N)$
	In middle	$O(N)$	$O(N)$	$O(N)$	$O(N)$
Linked List (doubly-linked) <u>Same as singly-linked single-ended!</u>	At front	$O(1)$	$O(1)$	$O(1)$	$O(1)$
	At end	$O(N)$	$O(N)$	$O(N)$	$O(N)$
	In middle	$O(N)$	$O(N)$	$O(N)$	$O(N)$

Complexity: Sorted List vs Simple List

Sorted list vs single-ended		Insert	Delete	Access	Find
Linked List (sorted) <u>Same as singly-linked single-ended!</u>	At front	$O(1)$	$O(1)$	$O(1)$	$O(1)$
	At end	$O(N)$	$O(N)$	$O(N)$	$O(N)$
	In middle	$O(N)$	$O(N)$	$O(N)$	$O(N)$

- Sorting will help only a little bit for find():
 - If the value *isn't* in the list ,we only need to search up to the point the number *should* be at, and then we can abort
 - Other lists are forced to search through all elements
- But it **doesn't** resolve poor scaling with larger N

Notes on Comparisons

- Arrays are unbeatable at $O(1)$ for access time
 - This makes them **indispensable**. Why? Because you often only add/remove an element *once*, but access an element *many times*
 - But if you only need to work on the **ends**, linked lists are better
- Single-ended lists are only good at working with **one end** (the head), double-ended lists are good at working with **both ends** (head and tail)
 - Otherwise, you might as well be using an array, UNLESS you need the ability to dynamically grow and shrink efficiently
- Doubly-*linked* lists don't help access/find time
 - How can you know it will be faster to go *backwards*?
 - Their real strength is application-specific back-and-forthing

SERIALIZATION


A cool way to save Objects

Persistence

- **Persistence:** The saving and loading application data so that it lasts between executions of the application
 - *i.e.*, the task of storing an application's state to disk and later loading it back to continue running
- We've already had a look at using text (and binary) formats to save data to disk and read it in later
 - Loading may involve parsing (text) saved data to translate it into a form that is suitable for manipulation by the application
 - Not all persistence requires parsing, but most do simply because storage is static whereas manipulation is dynamic

The Need for Serialization

- What if a ContainerClass (e.g.State) has a classfield which is an object of another class (aggregation)?
 - it does! Location class
- What if ContainerClass had an inheritance relationship that has classfields in the super class? e.g. Region in Country and State



COUNTRY:NAME=Australia:SHORTNAME=AUS:LANGUAGE=English:AREA=7692024:POPULATION=23401892:POPREF=Census2016
COUNTRY:NAME=Fiji:SHORTNAME=FIJ:LANGUAGE=English,Fijian,Hindi:AREA=18274:POPULATION=837271:POPREF=Census2016_Statoi
STATE:NAME=Western Australia:COUNTRY=Australia:SHORTNAME=WA:AREA=2529875:AREAUNIT=km2:POPULATION=2474410:POPREF
STATE:NAME=South Australia:COUNTRY=Australia:SHORTNAME=SA:AREA=983482:AREAUNIT=km2:POPULATION=1676653:POPREF=Ce
STATE:NAME=Victoria:COUNTRY=Australia:SHORTNAME=VIC:AREA=227416:AREAUNIT=km2:POPULATION=5926624:POPREF=Census20
STATE:NAME=New South Wales:COUNTRY=Australia:SHORTNAME=NSW:AREA=800642:AREAUNIT=km2:POPULATION=7480228:POPREF
STATE:NAME=Queensland:COUNTRY=Australia:SHORTNAME=QLD:AREA=1730648:AREAUNIT=km2:POPULATION=4703193:POPREF=Ce
STATE:NAME=Tasmania:COUNTRY=Australia:SHORTNAME=TAS:AREA=68401:AREAUNIT=km2:POPULATION=509965:POPREF=Census20
STATE:NAME=Northern Territory:COUNTRY=Australia:SHORTNAME=NT:AREA=1349129:AREAUNIT=km2:POPULATION=228833:POPREF
STATE:NAME=Australian Capital Territory:COUNTRY=Australia:SHORTNAME=ACT:AREA=2358:AREAUNIT=km2:POPULATION=396857:PO
STATE:NAME=Central Division:COUNTRY=Fiji:SHORTNAME=Central:AREA=4293:AREAUNIT=km2:POPULATION=342386:POPREF=Census
STATE:NAME=Western Division:COUNTRY=Fiji:SHORTNAME=Western:AREA=6360:AREAUNIT=km2:POPULATION=319611:POPREF=Cens
LOCATION:NAME=Perth:STATE=WA:COUNTRY=AUS:COORDS=-32.0024145,115.9097847:DESCRIPTION=State capital
LOCATION:NAME=Geraldton:STATE=WA:COUNTRY=AUS:COORDS=-28.7758955,114.5936913:DESCRIPTION=Coastal city in the Mid West
LOCATION:NAME=Albany:STATE=WA:COUNTRY=AUS:COORDS=-35.0258178,117.8754521:DESCRIPTION=Port city in the Great Southern reg
LOCATION:NAME=Port Hedland:STATE=WA:COUNTRY=AUS:COORDS=-20.3458226,118.5947316:DESCRIPTION=Major port in the North We
LOCATION:NAME=Melbourne Airport:STATE=VIC:COUNTRY=AUS:COORDS=-37.6690123,144.8388386:DESCRIPTION=Airport
LOCATION:NAME=Federation Square:STATE=VIC:COUNTRY=AUS:COORDS=-37.8177089,144.9668941:DESCRIPTION=Meeting place
LOCATION:NAME=Luna Park Melbourne:STATE=VIC:COUNTRY=AUS:COORDS=-37.8676749,144.9746723:DESCRIPTION=Fun park in St Kild

Serialization Issues

- Every object is a reference (pointer) – even Strings!
- Pointers let you jump around in memory (*i.e., non-contiguous*)
 - This is so useful that we simply cannot do without pointers!
- But you cannot save a pointer and expect it to be useful after re-loading it since the pointer says where the data *used to be* in the *old* run of the application!
 - When the application runs again, it won't place objects in the same place – every time a new object is created, it is given a different address
 - Furthermore, the object doesn't even exist yet – it is about to be loaded from disk!
- In other words, pointers cannot be made persistent

Saving Linked Lists to file

- To save to a file, we must put all the data together into a single block
- Only a few data structures are organised in blocks
 - Pretty much only arrays, and even then you need to add more info to define what the array's data type is
 - Each node of our linked list can be located in a completely different area of RAM
 - Remember, pointers cannot be made persistent
 - You cannot save a pointer and expect it to be useful after re-loading it since the pointer says where the data used to be in the old run of the application!
 - When the application runs again, it won't place objects in the same place – every time new is called, it returns a different address

How Does Serialization Work

- So how could it possibly work?
 - The idea is pack the data to be saved into a single contiguous block so that it can be written all together
 - Java and Python do this *automatically* via a language feature called reflection
- Reflection is the ability for an object to query the methods and fields of another class *at runtime*
 - This returns the names of fields, methods, parameters, return types, etc, and the ability to actually set/get the value of each field
 - Thus Java/Python can figure out what data fields an object has and pack them up for saving

Java Serialization

- Writing your own serialization code in Java is quite involved
 - The serialization buffer must be an array of bytes
 - Then you must copy every byte of an object into the serialization byte array
- This is why Java's serialization mechanism is so very useful
 - It does all the work of formatting and copying for you
 - From the programmer's viewpoint, it all comes down to the Serializable interface and a couple of Stream classes

Serializable Interface

- To let Java serialize your class, simply make it inherit from the Serializable interface *e.g.*,

```
public class ContainerClass implements Serializable
{
    ...
}
```

- That's it! **You do not have to write any code** – Java will determine the contents of your class and serialize it for you *entirely automatically* using reflection
- It is possible to override Java's serialization if needed
- Note that Serializable is an *empty* interface! In Java:

```
public interface Serializable { }
```

- *i.e.*, it is only used to mark a class as serializable for Java
- Don't write your own – import the Java one!

Object I/O Streams (Java)

- Any class that implements `Serializable` can then be written and read via two serialization-specific classes
 - `ObjectOutputStream` – for serializing an object and writing it out to an underlying stream
 - `ObjectInputStream` – for reading an underlying stream, deserializing the object and returning it
 - This creates the object and sets its fields to the serialized data
 - The two classes both know the serialization format
- When creating an Object stream, you must provide another stream – the underlying ‘physical’ stream
 - Usually `FileOutputStream` and `FileInputStream`

Object I/O Streams (Java)

- ObjectOutputStream will serialize the given object *and all objects referenced by that object*
 - This is done recursively, so that the entire object tree will be serialized
 - This supports the composition relationship: a car object consists of wheel objects, engine object and chassis object
 - So when serializing a car, we must serialize all contained objects
 - Objects not marked as Serializable won't be serialized
 - You can also add in custom serialization/deserialization steps, but this is only needed in special circumstances
- ObjectInputStream will deserialize all contained objs

Serializing an Object (Java)

```
private void save(ContainerClass objToSave, String filename)
{
    FileOutputStream fileStrm;
    ObjectOutputStream objStrm;

    try {
        fileStrm = new FileOutputStream(filename);
        objStrm = new ObjectOutputStream(fileStrm);
        objStrm.writeObject(objToSave);

        objStrm.close();

    }
    catch (Exception e) {
        throw new IllegalArgumentException("Unable to save object to file");
    }
}
```

← Underlying stream
← Object serialization stream
← Serialize and save to filename
 This will also save the ContainerClass' contained Location object
← Clean up

← should do more here

Deserializing an Object (Java)

```
private ContainerClass load(String filename) throws IllegalArgumentException
{
    FileInputStream fileStrm;
    ObjectInputStream objStrm;
    ContainerClass inObj;

    try {
        fileStrm = new FileInputStream(filename);           ← Underlying stream
        objStrm = new ObjectInputStream(fileStrm);          ← Object serialization stream
        inObj = (ContainerClass) objStrm.readObject();      ← Deserialize. Note the cast is needed
        objStrm.close();                                    ← Clean up
    }
    catch (ClassNotFoundException e) {
        System.out.println("Class ContainerClass not found" + e.getMessage());
    }
    catch (Exception e) {
        throw new IllegalArgumentException("Unable to load object from file");
    }
    return inObj;
}
```

Serialization in Python

- Python also uses reflection to convert object hierarchies into a bytestream for saving or transferring data
- The module for serialization is called `pickle`
- "`Pickling`" uses the `dump()` method to convert objects to bytestreams
- "`Unpickling`" uses the `load()` method to convert bytestreams to objects
- `dumps()` and `loads()` create and use strings instead of files
- **Note:** pickle doesn't protect your data, so only load data you trust

Pickling and Unpickling (Python)

Load Object

```
try:
    with open(placesDAT, "rb") as dataFile:
        inObject = pickle.load(dataFile)
except:
    print("Error: Object file does not exist!")
```

Save Object

```
print("Saving Object to File...")
try:
    with open(placesDAT, "wb") as dataFile:
        pickle.dump(myObject, dataFile)
except:
    print("Error: problem pickling object!")
```


Pickling and Unpickling (Python)

```
# Multiple objects
```

```
try:
    with open(placesDAT, "rb") as dataFile:
        countryObjects = pickle.load(dataFile)
        stateObjects = pickle.load(dataFile)
        locationObjects = pickle.load(dataFile)
except:
    print("Error: Object file does not exist!")
```

```
# Save Object
print("Saving Object to File...")
try:
    with open(placesDAT, "wb") as dataFile:
        pickle.dump(countryObjects, dataFile)
        pickle.dump(stateObjects, dataFile)
        pickle.dump(locationObjects, dataFile)
except:
    print("Error: problem pickling objects!")
```

XML and JSON Serialization

- The serialization approaches in previous slides use a format that is specific to the language(s)
 - Don't expect to be able to pass the object to another language and deserialize it, even if the classes are identical!
 - In fact, even backwards compatibility *in Java* has been broken
- An alternative that are the more portable XML and JSON formats
 - Both provide text-based, structured representations of objects
 - Java provides XMLEncoder and XMLDecoder for this, which work similarly to the Object<Input|Output>Stream classes
 - JAXB (Java Architecture for XML Binding) is a more sophisticated alternative for XML, but also more complex
 - JSON in python requires the json package and then is called similarly to pickle

ITERATORS

Next!!!!

Traversal

- Iterating over elements in a data set is a common task
 - e.g., to calculate the average of an array of ages:

```
public double calcAverageAge(double[] ages) {  
    double sum = 0.0;  
    for (int ii = 0; ii < ages.length; ii++) {  
        sum += ages[ii];  
    }  
    return sum / (double)ages.length;  
}
```

- An array's $O(1)$ time to access *any* element makes this efficient
- We could do the same with a LinkedList, but only if it provided some way to get at each element
 - **Solution 1:** Provide a peek(int index) method
 - **Solution 2:** Provide an *iterator* object

Inefficient LinkedList Traversal - Indexing

- The problem with Solution 1 is that it is inefficient
 - peek(int index) must start at the beginning of the list every time it is called
 - *i.e.*, it must traverse through the list from the beginning until it reaches the desired index
 - Thus a for loop to visit N elements will cause peek() to make N *traversals* of the list:

```
for (int ii = 0; ii < numAges; ii++)  
    sum += ages.peek(ii);
```

- Steps: $1 + 2 + 3 + 4 + \dots + N = N(N+1)/2$ steps $\approx N^2$
- Complexity: $O(N^2)$ (!!)
 - Truly awful scalability for a simple 'access each element' task

Efficient LinkedList Traversal - Iterators

- Solution 2 seeks to solve this inefficiency problem
 - Keep a reference to the *last visited* ('current') node
 - Sometimes called a cursor, since it indicates where you are
 - Every time the next node is requested, it only needs to make one hop forward
- Naïve approach: make `curr` a member of LinkedList
 - **Issue:** What if *two* loops access the list at the same time?
 - Commonly happens with nested for loops
 - But there's only *one* `curr` node in the LinkedList
 - Can't share it among many for loops

Iterators (Java)

- A LinkedList with a single ‘current’ cursor is limiting
 - Assumes only one ‘user/client’ of the LinkedList at a time
 - It would be better if every client had its own current cursor
- Iterators are designed to solve this problem
 - Each Java container (list) class has an iterator() method
 - Note that it’s not *getIterator*() - a bit inconsistent of Java!
 - Returns an Iterator (or more precisely, an object that inherits from the Java Iterator interface)
 - You don’t need to know the exact object type - knowing it is an Iterator is enough (the power of interfaces and polymorphism!)
 - Enumeration is an older Java interface for the same idea

Iterator Interface (Java)

- Quite simple; only three methods
 - hasNext() - queries if more items exist in the list
 - next() - move the cursor to the next item in the list
 - remove() - optional ability to remove the current item
 - throw UnsupportedOperationException if you don't support it
 - These just give standard names to a common task
- No prev() - inherit from interface ListIterator for that
- Only limit is that behaviour is undefined if element removed by one client while another client iterating
 - Depends on the underlying list being iterated over
 - Some can handle this scenario (e.g., linked list) others can't

Using an Iterator (Java)

- Using the iterator directly:

```
public void iterateOverList(LinkedList theList) {  
    MyClass c;  
    Iterator iter = theList.iterator();  
    while (iter.hasNext()) {  
        c = (MyClass)iter.next();           ← Get next item and cast from Object to MyClass  
        doSomething(c);  
    }  
}
```

- Using the ‘for-each’ looping structure:
 - Added to Java in 1.5 to simplify coding iterations
 - Requires list class to implement Iterable interface

```
public void iterateOverList(LinkedList theList) {  
    MyClass c;  
    for (Object o : theList) {  
        c = (MyClass)o;                     ← Cast to MyClass  
        doSomething(c);  
    }  
}
```

- Also handles nested loops without causing iteration errors

Writing Your Own Iterator (Java)

- Iterators are well-suited to being implemented by **private inner classes** - clients only need to know about Iterator

```
import java.util.*;

public class MyLinkedList implements Iterable {
    ...
    public Iterator iterator() {
        return new MyLinkedListIterator(this);
    }

    private class MyLinkedListIterator implements Iterator {
        private MyListNode iterNext;
        public MyLinkedListIterator(MyLinkedList theList) {
            iterNext = theList.head;
        }
        // Iterator interface implementation
        public boolean hasNext() { return (iterNext != null) }
        public Object next() {
            Object value;
            if (iterNext == null)
                value = null;
            else {
                value = iterNext.getValue();
                iterNext = iterNext.getNext();
            }
            return value;
        }
        public void remove() { throw new UnsupportedOperationException("Not supported"); }
    }
}
```

← so for-each loop can be used. Only defines iterator() method

← Return a new Iterator of internal type MyLinkedListIterator

← Hook the iterator to **this** MyLinkedList object

← Private class *inside* MyLinkedList

← Cursor (assuming MyListNode is the node class of MyLinkedList)

← **NOTE:** Able to access private field of MyLinkedList

← Get the value in the node

← Ready for subsequent calls to next()

Private Inner Classes (Java)

- Note that the iterator class was able to access the ‘head’ field of the list class
 - Even though ‘head’ is a private field!
 - Works because the inner class is also part of the list class
 - The inner class has **access to all private fields** of the outer class
 - This is a very useful property of inner classes that are exploited in many contexts
 - *e.g.*, event handling in GUIs, helper classes, and Iterators
 - Replaces the terrible concept of friend classes in C++
 - You could always pass the head instead of the list though

Iterators (Python)

- Need to define `__iter__()` and `__next__()` to build an iterator:
 - `__iter__` sets up your cursor to hold the current position in the collection
 - `__next__` defines how you move on to the next element
 - When you reach the end of the iteration, raise the **StopIteration** exception (e.g. hit the None at the end of a list)
- Once you have these, you can go through a list using a for loop:

```
for item in mylist:  
    print(item)
```

- `__iter__()` and `__next__()` are predefined functions, similar to `__init__()` and `__str__()`.
- A good resource for more information is :
https://www.w3schools.com/python/python_iterators.asp

Iterator Example (Python)

```
class ListNode():                                # incomplete sample code as guide
    def __init__(self, data):
        self._next = None                        # _ can be used to indicate
        self._data = data                       # "private" fields

class LinkedList():
    def __init__(self):
        self._root = None

    def insertFirst(self, value):
        if self._root == None:
            self._root = ListNode(value)
        else:
            newNode = ListNode(value)
            newNode._next = self._root
            self._root = newNode
```

Iterator Example (Python)

```
# class ListNode() continued
```

```
def __iter__(self):           # iter method creates iterator
    self._cur = self._root    # defines cursor for current
    return self               # returns reference to object

def __next__(self):          # next method defines traversal
    curval = None             # prepare return value
    if self._cur == None:     # check for end of list
        raise StopIteration   # Exception for end of iter
    else:
        curval = self._cur._data    # this return value
        self._cur = self._cur._next # move cursor along
    return curval              # return the value
```

```
# Problem: what if there's mutliple iterators on one list?
```

Iterator Example (Python)

```
# using the iterator
```

```
ll = LinkedList()  
ll.insertFirst("first")  
ll.insertFirst("second")  
ll.insertFirst("third")  
ll.insertFirst("fourth")
```

```
# creating list and getting next value
```

```
myiter = iter(ll)  
value = next(myiter)
```

```
# using a for loop to go through list
```

```
for value in ll:  
    print(value)
```

Another Iterator Approach (Python)

- Python lets you create a **generator** using the **yield** statement
- Maintains the current values of the method/function until next call
- **`__iter__()`** , **`__next__()`** and **StopIteration** are automatically implied
- Solves concurrent iterator issue as currNd is local to method

example of iterator using yield

```
def __iter__(self):  
    currNd = self.head  
    while currNd is not None:  
        yield currNd._value  
        currNd = currNd._nextNode
```


PRACTICAL 3

Aims:

- To create a general purpose Linked List class.
- To extend the linked list class with an iterator.
- To convert stack/queue to use the linked list.
- To save the linked list using serialization.

Next Week

- Trees

